

Problem Statement: -

A certain university wants to understand the relationship between students' SAT scores and their GPA. Build a Simple Linear Regression model with GPA as the target variable and record the RMSE and correlation coefficient values for different models.

In [1]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

df = pd.read_csv("D:\\360Digi\\Simple Resgression Ass\\SAT_GPA.csv")
df.isnull().sum()
```

Out[1]: SAT_Scores 0
GPA 0
dtype: int64

EDA

In [2]:

```
Eda = {"columns": df.columns,
      "mean": df.mean(),
      "median": df.median(),
      "stdrand deviation": df.std(),
      "variance": df.var(),
      "kurtosis": df.kurt()
      }

Eda
```

Out[2]: {'columns': Index(['SAT_Scores', 'GPA'], dtype='object'),
'mean': SAT_Scores 491.8100
GPA 2.8495
dtype: float64,
'median': SAT_Scores 480.5
GPA 2.8
dtype: float64,
'stdrand deviation': SAT_Scores 174.893834
GPA 0.541076
dtype: float64,
'variance': SAT_Scores 30587.853166
GPA 0.292764
dtype: float64,
'kurtosis': SAT_Scores -1.224188
GPA -1.040576
dtype: float64}

In [3]:

```
df.columns.values[0] = "SS"  
df.columns.values[1] = "Gpa"  
df.columns
```


Out[3]: Index(['SS', 'Gpa'], dtype='object')

In [4]:

```
plt.figure(figsize=(30, 30))  
sns.pairplot(df, hue='Gpa', height=3, diag_kind='hist')
```

Out[4]: <seaborn.axisgrid.PairGrid at 0x1fe2449c400>

<Figure size 2160x2160 with 0 Axes>




In [5]:

```
#yes or no count  
sns.catplot('Gpa', data=df, kind='count')
```

D:\anconda\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(


Out[5]: <seaborn.axisgrid.FacetGrid at 0x1fe24485d00>



In [6]:

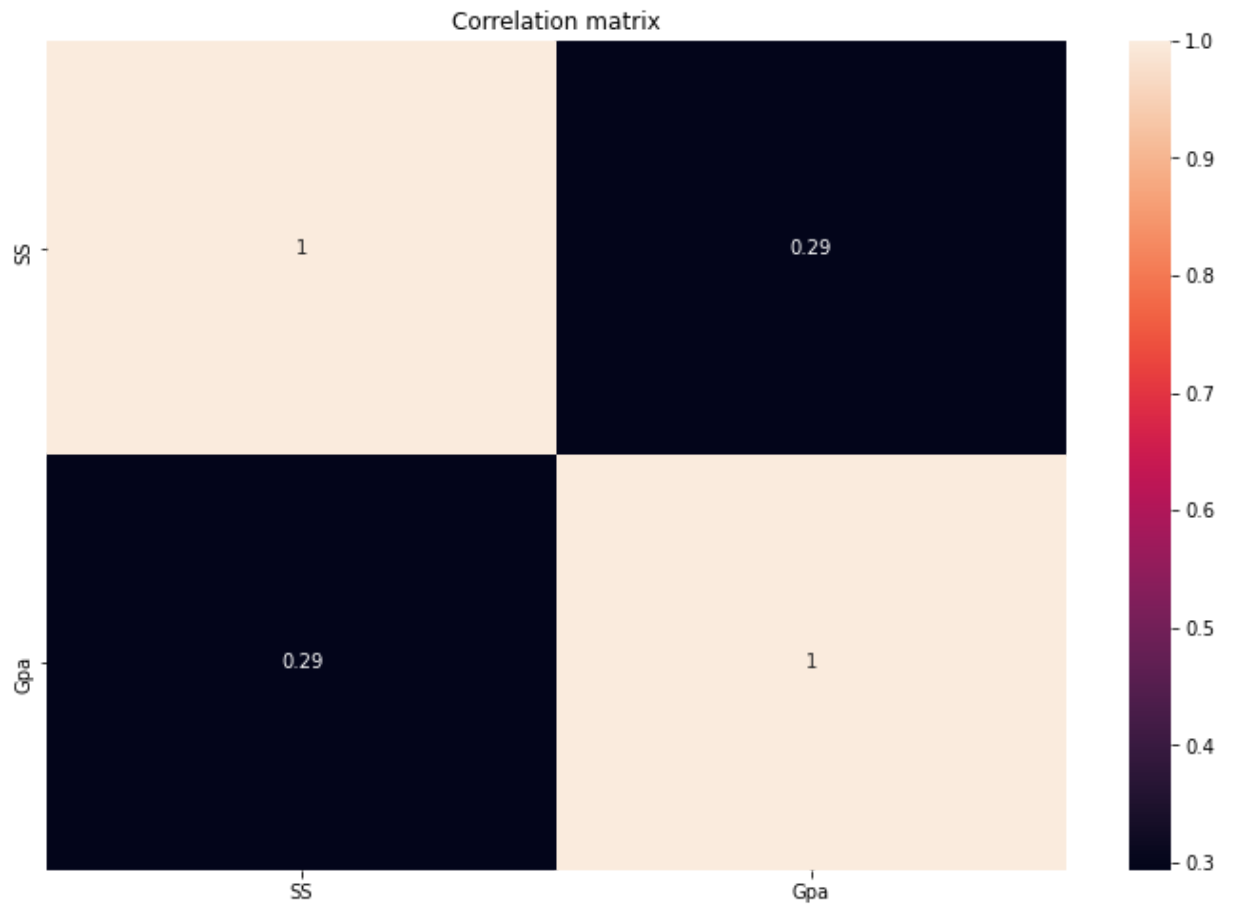
```
import matplotlib.pyplot as plt

plt.figure(figsize = (12, 8))
ax = sns.boxplot(data = df, orient = 'h', palette = 'Set2')
plt.title('Boxplot overview dataset')
plt.xlabel('values')
plt.xlim(-3, 500)
plt.show()
```



In [7]:


```
plt.figure(figsize = (12, 8))  
sns.heatmap(df.corr(), annot = True)  
plt.title('Correlation matrix')  
plt.show()
```



In [8]:

```
sns.pairplot(df)
```


Out[8]: <seaborn.axisgrid.PairGrid at 0x1fe28c4b8b0>



In [9]:

```
#Bar plot  
df['Gpa'].value_counts().plot.bar()
```

Out[9]: <AxesSubplot:>



In [10]:

```

'''
# Normalization function using z std. all are continuous data.
def std_func(i):
    x = (i-i.mean())/(i.std())
    return (x)

# Normalized data frame (considering the numerical part of data)
cal = std_func(df)
cal.describe()
'''

cal = df


#Data Modeling

#Graphical Representation
import matplotlib.pyplot as plt # mostly used for visualization purposes

#plt.bar(height = cal.ST, x = np.arange(1, 110, 1))
plt.hist(cal.SS) #histogram


```

Out[10]: (array([26., 17., 21., 19., 24., 17., 17., 20., 20., 19.]),
array([202. , 261.5, 321. , 380.5, 440. , 499.5, 559. , 618.5, 678. ,
737.5, 797.]),
<BarContainer object of 10 artists>)





```
In [11]: plt.boxplot(cal.SS) #boxplot
```

```
Out[11]: {'whiskers': [<matplotlib.lines.Line2D at 0x1fe2a306610>,  
  <matplotlib.lines.Line2D at 0x1fe2a306970>],  
  'caps': [<matplotlib.lines.Line2D at 0x1fe2a306cd0>,  
  <matplotlib.lines.Line2D at 0x1fe2a311070>],  
  'boxes': [<matplotlib.lines.Line2D at 0x1fe2a3062b0>],  
  'medians': [<matplotlib.lines.Line2D at 0x1fe2a3113d0>],  
  'fliers': [<matplotlib.lines.Line2D at 0x1fe2a311730>],  
  'means': []}
```




```
In [12]: plt.hist(cal.Gpa) #histogram
```

```
Out[12]: (array([23., 14., 34., 28., 25., 14., 9., 28., 11., 14.]),  
  array([2. , 2.19, 2.38, 2.57, 2.76, 2.95, 3.14, 3.33, 3.52, 3.71, 3.9 ]),  
  <BarContainer object of 10 artists>)
```




```
In [13]: plt.boxplot(cal.Gpa) #boxplot
```

```
Out[13]: {'whiskers': [<matplotlib.lines.Line2D at 0x1fe2a3dc7f0>,
<matplotlib.lines.Line2D at 0x1fe2a3dcb50>],
'caps': [<matplotlib.lines.Line2D at 0x1fe2a3dceb0>,
<matplotlib.lines.Line2D at 0x1fe2a3e8250>],
'boxes': [<matplotlib.lines.Line2D at 0x1fe2a3dc490>],
'medians': [<matplotlib.lines.Line2D at 0x1fe2a3e85b0>],
'fliers': [<matplotlib.lines.Line2D at 0x1fe2a3e8910>],
'means': []}
```



```
In [14]: # Scatter plot
plt.scatter(x = cal.SS, y = cal.Gpa, color = 'green')
```

```
Out[14]: <matplotlib.collections.PathCollection at 0x1fe2a43e400>
```



```
In [15]: # correlation
np.corrcoef(cal.SS, cal.Gpa)
```

```
Out[15]: array([[1.          , 0.29353828],
                [0.29353828, 1.          ]])
```

```
In [16]: # Covariance
# NumPy does not have a function to calculate the covariance between two variables
# Function for calculating a covariance matrix called cov()
# By default, the cov() function will calculate the unbiased or sample covariance matrix

cov_output = np.cov(cal.SS, cal.Gpa)[0, 1]
cov_output
```

```
Out[16]: 27.777793969849263
```

DATA MODELING

```
In [17]: # Import Library
import statsmodels.formula.api as smf

# Simple Linear Regression
model = smf.ols('Gpa ~ SS', data = cal).fit()
model.summary()
```

Out[17]: OLS Regression Results

Dep. Variable:	Gpa	R-squared:	0.086
Model:	OLS	Adj. R-squared:	0.082
Method:	Least Squares	F-statistic:	18.67
Date:	Sat, 19 Jun 2021	Prob (F-statistic):	2.46e-05
Time:	09:42:42	Log-Likelihood:	-151.44
No. Observations:	200	AIC:	306.9
Df Residuals:	198	BIC:	313.5
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	2.4029	0.110	21.908	0.000	2.187	2.619
SS	0.0009	0.000	4.321	0.000	0.000	0.001

Omnibus:	12.519	Durbin-Watson:	1.323
Prob(Omnibus):	0.002	Jarque-Bera (JB):	7.558
Skew:	0.317	Prob(JB):	0.0228
Kurtosis:	2.290	Cond. No.	1.56e+03

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


[2] The condition number is large, 1.56e+03. This might indicate that there are strong multicollinearity or other numerical problems.

In [18]:

```
pred1 = model.predict(pd.DataFrame(cal.SS))

# Regression Line
plt.scatter(cal.SS, cal.Gpa)
plt.plot(cal.SS, pred1, "r")
plt.legend(['Predicted line', 'Observed data'])
plt.show()

# Error calculation
res1 = cal.Gpa - pred1
res_sqr1 = res1 * res1
mse1 = np.mean(res_sqr1)
rmse1 = np.sqrt(mse1)
rmse1
```



Out[18]: 0.5159457227723684

In [19]:

```
##### Model building on Transformed Data
# Log Transformation
# x = log(waist); y = at

plt.scatter(x = np.log(cal.SS), y = cal.Gpa, color = 'brown')
np.corrcoef(np.log(cal.SS), cal.Gpa) #correlation

model2 = smf.ols('Gpa ~ np.log(SS)', data = cal).fit()
model2.summary()
```

Out[19]:

OLS Regression Results


Dep. Variable:	Gpa	R-squared:	0.077
Model:	OLS	Adj. R-squared:	0.072
Method:	Least Squares	F-statistic:	16.55
Date:	Sat, 19 Jun 2021	Prob (F-statistic):	6.85e-05
Time:	09:42:42	Log-Likelihood:	-152.42
No. Observations:	200	AIC:	308.8
Df Residuals:	198	BIC:	315.4
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	0.4796	0.584	0.822	0.412	-0.672	1.631
np.log(SS)	0.3868	0.095	4.068	0.000	0.199	0.574

Omnibus:	15.866	Durbin-Watson:	1.333
Prob(Omnibus):	0.000	Jarque-Bera (JB):	8.435
Skew:	0.320	Prob(JB):	0.0147
Kurtosis:	2.224	Cond. No.	99.8

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.




In [20]:

```
pred2 = model2.predict(pd.DataFrame(cal.SS))

# Regression Line
plt.scatter(np.log(cal.SS), cal.Gpa)
plt.plot(np.log(cal.SS), pred2, "r")
plt.legend(['Predicted line', 'Observed data'])
plt.show()

# Error calculation
res2 = cal.Gpa - pred2
res_sqr2 = res2 * res2
mse2 = np.mean(res_sqr2)
rmse2 = np.sqrt(mse2)
rmse2
```



Out[20]: 0.5184904101080668

In [21]:

```
#### Exponential transformation
# x = waist; y = log(at)
#cal.columns

plt.scatter(x = cal.SS, y = np.log(cal.Gpa), color = 'orange')
np.corrcoef(cal.SS, np.log(cal.Gpa)) #correlation

model3 = smf.ols('np.log(Gpa) ~ SS', data = cal).fit()
model3.summary()
```

Out[21]:

OLS Regression Results

Dep. Variable:	np.log(Gpa)	R-squared:	0.086
Model:	OLS	Adj. R-squared:	0.082
Method:	Least Squares	F-statistic:	18.75
Date:	Sat, 19 Jun 2021	Prob (F-statistic):	2.37e-05
Time:	09:42:43	Log-Likelihood:	58.615
No. Observations:	200	AIC:	-113.2
Df Residuals:	198	BIC:	-106.6
Df Model:	1		
Covariance Type:	nonrobust		


	coef	std err	t	P> t	[0.025	0.975]
Intercept	0.8727	0.038	22.745	0.000	0.797	0.948
SS	0.0003	7.35e-05	4.330	0.000	0.000	0.000

Omnibus:	11.046	Durbin-Watson:	1.375
Prob(Omnibus):	0.004	Jarque-Bera (JB):	4.816
Skew:	0.066	Prob(JB):	0.0900
Kurtosis:	2.251	Cond. No.	1.56e+03

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.56e+03. This might indicate that there are strong multicollinearity or other numerical problems.




In [22]:

```
pred3 = model3.predict(pd.DataFrame(cal.SS))
pred3_at = np.exp(pred3)
pred3_at

# Regression Line
plt.scatter(cal.SS, np.log(cal.Gpa))
plt.plot(cal.SS, pred3, "r")
plt.legend(['Predicted line', 'Observed data'])
plt.show()

# Error calculation
res3 = cal.Gpa - pred3_at
res_sqr3 = res3 * res3
mse3 = np.mean(res_sqr3)
rmse3 = np.sqrt(mse3)
rmse3
```



Out[22]: 0.5175875893834132

In [23]:

```
#### Polynomial transformation
# x = waist; x^2 = waist*waist; y = log(at)

model4 = smf.ols('np.log(Gpa) ~ SS + I(SS*SS)', data = cal).fit()
model4.summary()
```

Out[23]:

OLS Regression Results

Dep. Variable:	np.log(Gpa)	R-squared:	0.094
Model:	OLS	Adj. R-squared:	0.085
Method:	Least Squares	F-statistic:	10.23
Date:	Sat, 19 Jun 2021	Prob (F-statistic):	5.95e-05
Time:	09:42:43	Log-Likelihood:	59.448
No. Observations:	200	AIC:	-112.9
Df Residuals:	197	BIC:	-103.0
Df Model:	2		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	1.0056	0.110	9.112	0.000	0.788	1.223
SS	-0.0003	0.000	-0.607	0.545	-0.001	0.001
I(SS * SS)	6.142e-07	4.79e-07	1.284	0.201	-3.3e-07	1.56e-06

Omnibus:	8.598	Durbin-Watson:	1.357
Prob(Omnibus):	0.014	Jarque-Bera (JB):	4.118
Skew:	0.046	Prob(JB):	0.128
Kurtosis:	2.303	Cond. No.	2.79e+06

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.79e+06. This might indicate that there are strong multicollinearity or other numerical problems.

In [24]:

```


pred4 = model4.predict(pd.DataFrame(cal.SS))
pred4_at = np.exp(pred4)
pred4_at

# Regression Line
from sklearn.preprocessing import PolynomialFeatures
poly_reg = PolynomialFeatures(degree = 2)
X = cal.iloc[:, 0:1].values
X_poly = poly_reg.fit_transform(X)
# y = wcat.iloc[:, 1].values

plt.scatter(cal.SS, np.log(cal.Gpa))
plt.plot(X, pred4, color = 'red')
plt.legend(['Predicted line', 'Observed data'])
plt.show()

# Error calculation
res4 = cal.Gpa - pred4_at
res_sqr4 = res4 * res4
mse4 = np.mean(res_sqr4)
rmse4 = np.sqrt(mse4)
rmse4

```



Out[24]: 0.5144912487746158

In [25]:

```
# Choose the best model using RMSE
data = {"MODEL":pd.Series(["SLR", "Log model", "Exp model", "Poly model"]), "RMSE":pd.Series([0.515946, 0.518490, 0.517588, 0.514491])}
table_rmse = pd.DataFrame(data)
table_rmse
```

Out[25]:

	MODEL	RMSE
0	SLR	0.515946
1	Log model	0.518490
2	Exp model	0.517588
3	Poly model	0.514491

In [26]:

```
#####
# The best model

from sklearn.model_selection import train_test_split

train, test = train_test_split(cal, test_size = 0.3)

finalmodel = smf.ols('np.log(Gpa) ~ SS + I(SS*SS)', data = train).fit()
finalmodel.summary()
```

Out[26]:

OLS Regression Results

Dep. Variable:	np.log(Gpa)	R-squared:	0.106
Model:	OLS	Adj. R-squared:	0.093
Method:	Least Squares	F-statistic:	8.113
Date:	Sat, 19 Jun 2021	Prob (F-statistic):	0.000468
Time:	09:42:44	Log-Likelihood:	42.766
No. Observations:	140	AIC:	-79.53
Df Residuals:	137	BIC:	-70.71
Df Model:	2		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	0.9935	0.132	7.526	0.000	0.732	1.255
SS	-0.0003	0.001	-0.528	0.598	-0.001	0.001
I(SS * SS)	6.43e-07	5.66e-07	1.135	0.258	-4.77e-07	1.76e-06

Omnibus:	3.861	Durbin-Watson:	2.230
Prob(Omnibus):	0.145	Jarque-Bera (JB):	2.566
Skew:	0.143	Prob(JB):	0.277
Kurtosis:	2.402	Cond. No.	2.78e+06

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.78e+06. This might indicate that there are strong multicollinearity or other numerical problems.

In [27]:

```
# Predict on test data
test_pred = finalmodel.predict(pd.DataFrame(test))
pred_test_AT = np.exp(test_pred)
pred_test_AT

# Model Evaluation on Test data
test_res = test.Gpa - pred_test_AT
test_sqr = test_res * test_res
test_mse = np.mean(test_sqr)
test_rmse = np.sqrt(test_mse)
test_rmse
```

Out[27]: 0.537987761257369

In [28]:

```
# Prediction on train data
train_pred = finalmodel.predict(pd.DataFrame(train))
pred_train_AT = np.exp(train_pred)
pred_train_AT

# Model Evaluation on train data
train_res = train.Gpa - pred_train_AT
train_sqr = train_res * train_res
train_mse = np.mean(train_sqr)
train_rmse = np.sqrt(train_mse)
train_rmse
```

Out[28]: 0.5082608011769393

Summary

Model having highest R-Squared value is better i.e. (model=0.897 is not better than model1=0.960). There has good relationship>0.85

RMSE- lower the RMSE indicate better fit. RMSE is a good measure of how accuracy the model predict the response. In Linear regression RMSE value between 0.2 to 0.5

But in final model training and testing we choose Polynomial transformation $\ln(\text{Gpa}) \sim \text{SS} + \ln(\text{SS} \times \text{SS})$ because the training rmse was show good result in Polynomial transformation rather than SLR,Log.

In []:

